

PhaseFieldPet: An Open-Source Phase-Field Modeling Software for Heterogeneous Architectures

Shore Salle Chota^{1,4}, Martin Reder^{1,2}, Daniel Schneider^{1,2}, Harald Koestler⁴, and Britta Nestler^{1,2,3}

¹ Institute for Applied Materials (IAM), Karlsruhe Institute of Technology, Karlsruhe, 76131, Germany ² Institute of Digital Materials Science (IDM), Karlsruhe University of Applied Sciences, Karlsruhe, 76133, Germany ³ Institute of Nanotechnology (INT), Karlsruhe Institute of Technology, Karlsruhe, 76131, Germany ⁴ Chair for System Simulation, Friedrich-Alexander-Universität Erlangen-Nürnberg, Erlangen, Germany ¶ Corresponding author

DOI: 10.xxxxxx/draft

Software

- Review
- Repository
- Archive

Editor: Open Journals

Reviewers:

- @openjournals

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License (CC BY 4.0).

Summary

Phase field method has emerged as a powerful computational tool for simulating complex, evolving interfaces at mesoscale in materials science and engineering, fluid dynamics, cell migration and other fields. However, achieving scalability and efficiency for parallel simulations of multicomponent multiphase systems across diverse hardware architectures remains a challenge. This paper presents PhaseFieldPet, an open-source, message passing interface (mpi) based software package for large-scale phase field simulations, specifically designed to leverage heterogeneous architectures such as CPUs and GPUs. PhaseFieldPet is built upon the Portable, Extensible Toolkit for Scientific Computation (PETSc), to efficiently handle large-scale simulations on high-performance computing platforms. The software's modular design facilitates and easily integrates various phase field models such as Multiphase-Field and Multi-Order Parameter models with various choice of gradient and potential energy contributions. Performance benchmarks demonstrate the software's capability to handle simulations involving from small to millions of degrees of freedom with excellent scalability.

The Phase Field Equation

The phase-field approach represents surfaces and interfaces implicitly using continuous scalar fields (order parameter) $\phi_\alpha(\mathbf{r}, t)$, $\alpha \in \{1, 2, \dots, N\}$, which are constant in bulk phases and transition smoothly—but sharply—across a diffuse boundary. The order parameters $\phi_\alpha(\mathbf{r}, t)$ represent different grains or phases such as solid, liquid, gas and their state like crystallographic orientation, polarization, volume fraction. For example, in a solid-solid phase transition, a solid phase can be represented by N order parameters $\phi_\alpha(\mathbf{r}, t)$ based on N crystallographic orientations or grains.

Microstructure evolution, and hence evolution of order parameter $\phi_\alpha(\mathbf{r}, t)$ can be obtained from functionals of entropy, or free energy, or grand potential (Hötzer et al., 2018). Following an energy approach, one can write the total free energy functional of the system as

$$\mathcal{F}(\phi, \nabla\phi, \dots) = \int_V f dV = \int_V f_{\text{grad}}(\phi, \nabla\phi) + f_{\text{pot}}(\phi) + f_{\text{bulk}}(\phi, \dots) dV.$$

The first two terms contribute to interfacial energy, and exist only at points, where multiple order parameters are non-zero, and thus interfaces occur. These terms are responsible to keep the diffuse interface at finite width with interplay of gradient energy density $f_{\text{grad}}(\phi, \nabla\phi)$ which diffuses the interfaces, while the potential term $f_{\text{pot}}(\phi)$ counter acts it. In equilibrium

these terms are equal. The bulk contribution $f_{\text{bulk}}(\phi, \dots)$ can be of various type depending on the problem at hand such as chemical, thermal, mechanical, electrical, magnetic and etc.

There exists various formulations to $f_{\text{grad}}(\phi, \nabla\phi)$ and $f_{\text{pot}}(\phi)$ by many scholars in the Phase field community (Daubner et al., 2023). As an example, (Nestler et al., 2005) formulate these terms as

$$f_{\text{grad}}(\phi, \nabla\phi) = \varepsilon \sum_{\alpha} \sum_{\beta > \alpha} \gamma_{\alpha\beta} |\phi_{\alpha} \nabla\phi_{\beta} - \phi_{\beta} \nabla\phi_{\alpha}|^2,$$

$$f_{\text{pot}}(\phi) = \frac{16}{\varepsilon \pi^2} \sum_{\alpha} \sum_{\beta > \alpha} \gamma_{\alpha\beta} \phi_{\alpha} \phi_{\beta} + \frac{1}{\varepsilon} \sum_{\alpha} \sum_{\beta > \alpha} \sum_{\delta > \beta} \gamma_{\alpha\beta\delta} \phi_{\alpha} \phi_{\beta} \phi_{\delta}.$$

See Daubner et al. (2023) table B.4 for other formulations of these terms and their explanations currently considered in PhaseFieldPet.

Following total energy minimization principle of the system, the phase field evolution equations for each $\phi_{\alpha}(\mathbf{r}, t)$ are in general derived by variational derivatives of the functional $\mathcal{F}(\phi, \nabla\phi, \dots)$ and are given by Allen-Cahn or time-dependent Ginzburg–Landau equations for each order parameter by

$$\frac{\partial\phi_{\alpha}}{\partial t} = -L \frac{\delta\mathcal{F}}{\delta\phi_{\alpha}} = -L \left(\frac{\partial f}{\partial\phi_{\alpha}} - \nabla \cdot \frac{\partial f}{\partial\nabla\phi_{\alpha}} \right),$$

where L is kinetic coefficient. The resulting set of nonlinear parabolic differential equations define the multi-order parameter (MOP) phase-field model and are selected in PhaseFieldPet via option `pfe_mop`.

Multiphase-field models restrict the phase fields such that $\phi_{\alpha}(\mathbf{r}, t) \in [0, 1]$, $\sum_{\alpha} \phi_{\alpha} = 1$. (Nestler et al., 2005) introduced a Lagrange multiplier λ yielding Allen-Cahn type Phase field equation

$$\frac{\partial\phi_{\alpha}}{\partial t} = -L \left(\frac{\partial f}{\partial\phi_{\alpha}} - \nabla \cdot \frac{\partial f}{\partial\nabla\phi_{\alpha}} \right) - \lambda.$$

This is Lagrangian based Multiphase-field model (`mpfl`), and is chosen in PhaseFieldPet by `pfe_mpfl`.

(Steinbach & Pezzolla, 1999) rewrote the phasefield evolution equations by the sum of binary interactions

$$\frac{\partial\phi_{\alpha}}{\partial t} = -\frac{1}{\tilde{N}\epsilon} \sum_{\beta \neq \alpha}^{\tilde{N}} M_{\alpha\beta} \left(\frac{\delta\mathcal{F}}{\delta\phi_{\alpha}} - \frac{\delta\mathcal{F}}{\delta\phi_{\beta}} \right),$$

where $M_{\alpha\beta}$ is a mobility matrix. This Multiphase-field model (`mpf`), is chosen in PhaseFieldPet via `pfe_mpf`.

We refer interested readers to Daubner et al. (2023), Moelans et al. (2008) and Chapter seven of the book by Provatas & Elder (2010) for a detailed overview of various phase-field formulations and associated evolution equations.

Statement of need

For the past couple of decades, phase-field software has been being developed and used with in house codes, and Open source phase field software started to be available from 2007 (Hong & Viswanathan, 2020). Many existing open source softwares are limited to one or two spatial dimensions, focus on binary systems, use only one type time step solver (usually explicit time stepping), work only on one CPU core (serial code) or are not capable of using heterogeneous compute resources available such as GPUs for compute and energy efficiency. Notable large scale, distributed computing capable open source phase-field software that mainly target CPUs include: The open source Multiphysics Object Oriented Simulation Environment (MOOSE) (Schwen et al., 2017) - which is a powerful toolset for implementing phase-field models using the finite element method (FEM), PRISMS-PF - massively parallel finite element code for conducting phase field (DeWitt et al., 2020), OpenPhase (Tegeler et al., 2017) uses finite difference method (FDM) for spatial discretization, an explicit time stepping algorithm, MicroSim (Dutta et al., 2025). Among proprietary, distributed machines capable software is a Parallel Algorithms for Crystal Evolution in 3D (PACE3D) (Hötzer et al., 2018) is a multiphase-field software that uses explicit time stepping along with finite difference spatial discretizations. Table 1 below gives a comparison of selected state of the art software for Allen-Cahn (and variations thereof) type phase-field model solvers with online tutorial available, able to run on distributed - large scale hardware architectures.

Software	Various Time Step Solver	3D capability	GPU capability	Phase Field Model	Spatial Discretization	Remark
PRISMS-PF	No	Yes	No	mop	FEM	Built on Deal.II
MicroSim	No	Yes*	Yes*	mpfl*, mpf	FDM	* mpfl is not 3D and GPU capable yet
OpenPhase	No	Yes	No	mpf	FDM	GPU capability ongoing
PACE3D	No	Yes	No	mpfl, mpf	FDM	Proprietary
MOOSE	Yes	Yes	Yes**	mop, mpfl	FEM	**Possible for Nonlinear solver (via PETSc)
PhaseFieldPet	Yes	Yes	Yes	mop, mpfl, mpf	FDM	FEM possible

Table 1: MPI capable phase-field software.

PhaseFieldPet is a finite difference method based software built on top of Time Step (TS) solver from PETSc (Abhyankar et al., 2018). It is based on the previous work (Daubner et al., 2023) including all the different model formulations compared therein. PhaseFieldPet extends the numerical solutions to 3D, an arbitrary amount of N phases and includes various bulk driving forces according to Hoffrogge et al. (2025). It fills the aforementioned gaps in existing software by combining the following features:

1. Enabling of multiphase simulations in 1D / 2D / 3D.
2. Decoupling the numerical solution methods from the physical modeling such that one can choose various solution methods without restricting to one time step solver (i.e. one can use methods like semi implicit, implicit time stepping algorithms, various underlying nonlinear solver, linear solvers and preconditioners, etc) based on composability features of PETSc (Balay et al., 2024). This makes it easier for newcomers to the phase-field community and advanced users alike. The set of phase-field equations is formulated using the Implicit Explicit (IMEX) scheme such that either the whole phase-field equation is treated implicitly or the stiff part, allowing longer time steps compared to the explicit methods which require small time steps for stability.
3. Executable on single core, multicore to multi node High Performance Computing clusters/supercomputers coupled with accelerators such as GPUs (Mills et al., 2021). GPUs

are increasingly available for computing purposes with thousands of compute cores and small energy usage. Compute power is leveraged using PhaseFieldPet from one code base for both CPUs and GPUs.

4. Flexibility to easily switch between various phase-field models and energy contributions at run time.

Usage

To use PhaseFieldPet, PETSc (Balay et al., 2024) needs to be installed, compiled using

```
make PhaseFieldPet
```

run the executable generated with default solver settings (E.g. with 4 mpi process) by

```
mpiexec -n 4 PhaseFieldPet
```

By default PhaseFieldPet simulates the benchmark case introduced by Daubner et al. (2023), considering a stationary triple junction problem. The default model configuration is the usage of dot gradient term (grad_dot), the well potential of Toth (pot_toth) and a Lagrange multiplier based multiphase-field formulation (pfe_mpfl). The default time stepping solver being adaptive Runge Kutta Implicit-Explicit (IMEX) method, where the stiff part of the equation is treated implicitly.

To simulate with non default combinations, for instance to solve phase-field equations with gradient and potential energy terms of Nestler et al. (2005) and with $f_{\text{bulk}}(\phi, \dots) = 0$, we give the corresponding options to the executable as

```
mpiexec -n 4 PhaseFieldPet -grad_weighted -pot_nestler -simplex
```

In this configuration, a weighted (generalized) gradient energy formulation (grad_weighted), the obstacle potential (pot_nestler) and a Gibbs simplex constraint (simplex) to constrain each phase field $\phi_\alpha \geq 0$, $\sum_\alpha \phi_\alpha = 1$ at each point in the simulation domain is composed. Setting up of different phase-field equations (pfe) is enabled with the option like pfe_mop (Multi-order parameter model) not restricting the order parameters, pfe_mpf (Non lagrangian based Multi-Phase field Equation) (Tegeler et al., 2017) along with other gradient and potential energy contributions. For details of usage not mentioned here, including your own energy contributions, see the associated Github page to this paper.

Example Performance Result

Strong scalability of PhaseFieldPet is achieved for simulations of static triple junctions using second order, adaptive Backward Euler (fully implicit) time step solver by increasing grid points along x and y direction using

```
mpiexec -n # PhaseFieldPet -simplex -ts_type bdf -da_grid_x 256 -da_grid_y 256
```

Figure 1 shows the result on Meggie cluster at NHR@FAU obtained by running up to 80 mpi processes on four compute nodes, where each node has two Intel Xeon E5-2630v4 “Broadwell” chips (10 cores per chip) running at 2.2 GHz with 25 MB Shared Cache per chip and 64 GB of RAM. The result indicates excellent agreement with ideal expectations that the log-log plot is a straight line with slope of -1.




Figure 1: Scalability plot: Log-Log Plot of Execution Time vs MPI Processes.

Conclusions

PhaseFieldPet provides users with flexible methods to solve multiphase-field equations, along with various energy contributions on heterogeneous hardware architectures including CPUs and GPUs. More specifically, a user can include specific gradient, potential, bulk driving energy terms and choose the type of the phase field equation, and type of numerical algorithm to use in order to solve the differential equation (along with the choice of the underlying nonlinear equation solver, linear equation solver, preconditioners) and etc. Inline with PACE3D software (Hötzer et al., 2018) and its extension, the future version of PhaseFieldPet will include various other modules corresponding to different applications of general multiphysics multiphase-field methods.

Acknowledgement

We acknowledge discussions with Dr. Simon Daubner during the genesis of this project. The authors are grateful for the scientific support and HPC resources provided by the Erlangen National High Performance Computing Center (NHR@FAU) of the Friedrich-Alexander-Universität Erlangen-Nürnberg (FAU). The hardware is funded by the German Research Foundation (DFG).

Funding

The author SSC thanks the NHR-Verein e.V for supporting this work/project within the NHR Graduate School of National High Performance Computing (NHR). Furthermore, DS and BN acknowledge support from the Helmholtz association within programme MTET, no.38.04.04.

References

- Abhyankar, S., Brown, J., Constantinescu, E. M., Debojyoti Ghosh, Smith, B. F., & Zhang, H. (2018). PETSc/TS: A modern scalable ODE/DAE solver library. In *arXiv e-prints*. <https://arxiv.org/abs/1806.01437>
- Balay, S., Abhyankar, S., Adams, M. F., Benson, S., Brown, J., Brune, P., Buschelman, K., Constantinescu, E. M., Dalcin, L., Dener, A., Eijkhout, V., Faibussowitsch, J., Gropp, W.

- D., Hapla, V., Isaac, T., Jolivet, P., Karpeev, D., Kaushik, D., Knepley, M. G., ... Zhang, J. (2024). *PETSc Web page*. <https://petsc.org/>
- Daubner, S., Hoffrogge, P. W., Minar, M., & Nestler, B. (2023). Triple junction benchmark for multiphase-field and multi-order parameter models. *Computational Materials Science*. <https://doi.org/10.1016/j.commatsci.2022.111995>
- DeWitt, S., Rudraraju, S., D., M., Andrews, W. B., & Thornton, K. (2020). PRISMS-PF: A general framework for phase-field modeling with a matrix-free finite element method. *Npj Computational Materials*. <https://doi.org/10.1038/s41524-020-0298-5>
- Dutta, T., Mohan, D., Shenoy, S., Attar, N., Kalokhe, A., Sagar, A., Bhure, S., Swaroop S. Pradhan, SS., Praharaj, J., Mridha, S., Kushwaha, A., Shah, V., Gururajan, M. P., Venkatesh Sheno, V., Phanikumar, G., Bhattacharyya, S., & Choudhury, A. (2025). MicroSim: A high-performance phase-field solver based on CPU and GPU implementations. *Computational Materials Science*, 246, 113438. <https://doi.org/10.1016/j.commatsci.2024.113438>
- Hoffrogge, P. W., Daubner, S., Schneider, D., Nestler, B., Zhou, B., & Eiken, J. (2025). Triple junction benchmark for multiphase-field models combining capillary and bulk driving forces. *Modelling Simul. Mater. Sci. Eng.* <https://doi.org/10.1088/1361-651X/ad8d6f>
- Hong, Z., & Viswanathan, V. (2020). Open-Sourcing Phase-Field Simulations for Accelerating Energy Materials Design and Optimization. *ACS Energy Letters*. <https://doi.org/10.1021/acsenergylett.0c01904>
- Hötzer, J., Reiter, A., Hierl, H., Steinmetz, P., Selzer, M., & Nestler, B. (2018). The parallel multi-physics phase-field framework Pace3D. *Journal of Computational Science*. [10.1016/j.jocs.2018.02.011](https://doi.org/10.1016/j.jocs.2018.02.011)
- Mills, R. T., Adams, M. F., Balay, S., Brown, J., & Dener, A. (2021). Toward performance-portable PETSc for GPU-based exascale systems. *Parallel Computing*, 108, 102831. <https://doi.org/10.1016/j.parco.2021.102831>
- Moelans, N., Blanpain, B., & Wollants, P. (2008). An introduction to phase-field modeling of microstructure evolution. *Calphad*, 134, 268–294. <https://doi.org/10.1016/j.calphad.2007.11.003>
- Nestler, B., Garcke, H., & Stinner, B. (2005). Multicomponent alloy solidification: Phase-field modeling and simulations. *Physical Review E*. <https://doi.org/10.1016/j.commatsci.2022.111995>
- Provatas, N., & Elder, K. (2010). *Phase-field methods in materials science and engineering*. Wiley-VCH Verlag GmbH & Co. KGa. <https://doi.org/10.1002/9783527631520>
- Schwen, D., Aagesen, L. K., Peterson, J. W., & Tonks, M. R. (2017). Rapid multiphase-field model development using a modular free energy based approach with automatic differentiation in MOOSE/MARMOT. *Computational Materials Science*, 132, 36–45. <https://doi.org/10.1016/j.commatsci.2017.02.017>
- Steinbach, I., & Pezzolla, F. (1999). Generalized field method for multiphase transformations using interface fields. *Physica D*, 134, 385–393. [https://doi.org/10.1016/S0167-2789\(99\)00129-3](https://doi.org/10.1016/S0167-2789(99)00129-3)
- Tegeler, M., Shchyglo, O., Kamachali, R. D., Monas, A., Steinbach, I., & Sutmann, G. (2017). Parallel multiphase field simulations with OpenPhase. *Computer Physics Communications*, 215, 173–187. <https://doi.org/10.1016/j.cpc.2017.01.023>